



TECHNISCHE UNIVERSITÄT
BERGAKADEMIE FREIBERG

Die Ressourcenuniversität. Seit 1765.

Prof. Dr. Björn Sprungk

Faculty of Mathematics and Computer Science

Institute of Stochastics

Mathematics of machine learning

5. Training

Winter term 2024/25

5. Training

Content

5.1 Validation and testing

5.2 Gradient descent

5.3 Stochastic gradient descent

5.4 Outlook on nonsmooth optimization

5. Training

What's it about?

We will look at some aspects of optimization in machine learning:

1. How should you use the available data also for validation and testing?
2. How can empirical risk minimization be performed in practice using numerical optimization?

⇒ Gradient descent

3. How is learning done in practice with large or even huge amounts of data?

⇒ Stochastic gradient descent!

5.1 Training, validation, testing



- One should use the available data for different purposes:
 1. **Training:** Determine hypothesis h_s via numerical computation of a learning rule $h_s \approx A(s)$, e.g. ERM rule $A = \text{ERM}_{\mathcal{H}}$
 2. **Validation or model selection:** Choose the most appropriate hypothesis class \mathcal{H} or hyperparameter such as regularization parameter λ to train
 3. **Testing:** Estimate the generalization error of the learned hypothesis
- To this end, we divide the data $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_m, y_m)$ **randomly** into a **training sample s of m_S points**, a **validation sample v of m_V points**, and a **test set t of m_T points**
- **Rule of thumb:** $m_V = m_T$ and $m_S \in [0.6m, 0.8m]$ – i.e. 60% to 80% percent of data for training and 20% and 10% for validation and testing.

Validation

By validation we identify the best **hyperparameters** for the learning task, e.g., the **regularization parameter** λ or a **complexity parameter** p of the hypothesis class.

Example 5.1:

- Given $m = 10$ data pairs (x_i, y_i) we want to fit or learn a polynomial $h_p(x) = \sum_{j=0}^p w_j x^j$
- Hypotheses $h_{s,p}$ determined by ERM for $p = 2$, $p = 3$ and $p = 10$:





- Which of these leads to the smallest generalization error $\mathcal{R}_\mu(h_{s,p})$?

We compare different hypotheses $h_{s,1}, \dots, h_{s,J}$ learned by (regularized) ERM from J classes or models $\mathcal{H}_1, \dots, \mathcal{H}_J$ by their **validation error** based on the **validation dataset** $v = ((\mathbf{x}_1, y_1), \dots, (\mathbf{x}_{m_V}, y_{m_V}))$:

$$\mathcal{R}_v(h_{s,j}) = \frac{1}{m_V} \sum_{i=1}^{m_V} \ell(h_{s,j}(\mathbf{x}_i), y_i).$$

Example 5.1:

For the given example of polynomial fitting we obtain with $m_V = 16$ additional validation points:



Source: "Understanding Machine Learning" (2014)

\Rightarrow We choose $h_{s,p}$ for $p = d = 3$ since there the **validation error** is smallest

Quality of the validation error

- The **validation error** $\mathcal{R}_v(h_s)$ shall estimate the expected risk $\mathcal{R}_\mu(h_s)$ analogously to \mathcal{R}_s
- Can we bound the difference $|\mathcal{R}_\mu(h_s) - \mathcal{R}_v(h_s)|$?
- We can apply **Hoeffding's inequality** and obtain for bounded loss $\ell \leq c$:

$$\mathbb{P}_{\mu^{m_V}} (|\mathcal{R}_V(h_s) - \mathcal{R}_\mu(h_s)| > \epsilon) \leq 2 \exp(-2m_V \epsilon^2 / c^2)$$

- For $\epsilon = \sqrt{\frac{c^2 \log(2/\delta)}{2m_V}}$ we get $\mathbb{P} (|\mathcal{R}_V(h_s) - \mathcal{R}_\mu(h_s)| \leq \epsilon) \geq 1 - \delta$
- By **Bonferroni's correction** we obtain for J hypotheses to be compared

$$\mathbb{P}_{\mu^{m_V}} \left(|\mathcal{R}_V(h_{s,j}) - \mathcal{R}_\mu(h_{s,j})| \leq \sqrt{\frac{c^2 \log(2J/\delta)}{2m_V}} \quad \forall j = 1, \dots, J \right) \geq 1 - \delta$$

Cross validation

1. We divide the m data points into K blocks s_k , $k = 1, \dots, K$, of $\frac{m}{K}$ points – assuming $\frac{m}{K} \in \mathbb{N}$
2. For each of blocks s_k compute for the J hyperparameters (e.g., for classes $\mathcal{H}_1, \dots, \mathcal{H}_J$ or regularization parameters $\lambda_1, \dots, \lambda_J$) the hypotheses

$$h_{k,j} = A_j((s_k^*)), \quad s_k^* = (s_1, \dots, s_{k-1}, s_{k+1}, \dots, s_K),$$

where A_j denotes the learning algorithm with respect to the j -th hyperparameter, and calculate

$$\text{err}_{k,j} := \mathcal{R}_{s_k}(h_{k,j})$$

3. The cross validation error for the j -th learning model or the j -th hyperparameter is then given by

$$\text{err}_j^{\text{cv}} := \frac{1}{K} \sum_{k=1}^K \text{err}_{k,j} = \frac{1}{K} \sum_{k=1}^K \mathcal{R}_{s_k}(h_{k,j})$$

4. We then choose $j^* \in \operatorname{argmin}_j \text{err}_j^{\text{cv}}$ and compute $h_s = A_{j^*}(s)$

- Cross validation is particularly applied if there is not enough data to split it into a training and a validation set
- With the Hoeffding inequality one can control again for each j and k $|\mathcal{R}_\mu(h_{k,j}) - \mathcal{R}_{s_k}(h_{k,j})|$ but no longer $|\mathcal{R}_\mu(h_s) - \text{err}_{j^*}^{\text{CV}}|$ due to the dependencies of the errors $\text{err}_{k,j}$.
- In practice, cross validation usually works very well, but in theory it can also lead to a wrong choice.
- The case $K = m$ is also called [Leave-one-out \(LOO\) cross validation](#).

- Once one has decided for learned hypothesis h_s , a **test dataset** $t = ((\mathbf{x}_1, y_1), \dots, (\mathbf{x}_{m_T}, y_{m_T}))$ can be used to determine the **test error**

$$\mathcal{R}_T(h_s) = \frac{1}{m_T} \sum_{i=1}^{m_T} \ell(h_s(\mathbf{x}_i), y_i),$$

as an estimate of the generalization error $\mathcal{R}_\mu(h_s)$.

- Again, using the Hoeffding inequality for restricted loss functions we have

$$\mathbb{P}_{\mu^{m_T}} \left(|\mathcal{R}_T(h_s) - \mathcal{R}_\mu(h_s)| \leq \sqrt{\frac{c^2 \log(2/\delta)}{2m_T}} \right) \geq 1 - \delta.$$

- If a validation error has already been calculated for h_s based on a validation data set v , then this can also be used as an estimate.

Training via numerical optimization

- We will deal with numerical optimization for computing the outcome of the (regularized) ERM rule

$$h_s = A(s) = \operatorname{argmin}_{h \in \mathcal{H}} \mathcal{R}_s(h) + \lambda R(h)$$

consisting of an empirical risk $\mathcal{R}_s(h)$ and, if necessary, a regularization $R(h)$ in the remainder of this chapter.

- In practice, the minimizer of $\mathcal{R}_s + \lambda R$ is calculated via numerical optimization methods.
- There are nowadays many such methods adapted to machine learning like [AdaGrad](#) or [ADAM](#).
- We restrict ourselves here to the two classical and simplest ones, gradient decent and **stochastic gradient descent**.

Parametrized hypothesis classes

- We assume subsequently that the hypothesis classes under consideration $\mathcal{H} \subseteq \mathcal{Y}^{\mathcal{X}}$ are **parameterized**.
- That is, there exists a **parameter set** $\mathcal{W} \subseteq \mathbb{R}^p$, $p \in \mathbb{N}$, and each hypothesis $h \in \mathcal{H}$ corresponds to a parameter(vector) $\mathbf{w} \in \mathcal{W}$: $h = h_{\mathbf{w}}$.
- The mapping $\mathbf{w} \mapsto h_{\mathbf{w}}$ does not need to be injective, e.g., for linear hypotheses $h_{(\mathbf{w},b)} \in \mathcal{L}_d$ we have $h_{\lambda(\mathbf{w},b)} \equiv h_{(\mathbf{w},b)}$ for all $\lambda > 0$.
- Furthermore, we assume now a **loss function** $\ell: \mathcal{W} \times \mathcal{X} \times \mathcal{Y} \rightarrow [0, \infty)$ stated on the parameter set w. r. t which we have want to compute

$$\mathbf{w}_S \in \operatorname{argmin}_{\mathbf{w} \in \mathcal{W}} \mathcal{R}_s(\mathbf{w}) + \lambda R(\mathbf{w}), \quad \mathcal{R}_s(\mathbf{w}) := \frac{1}{m} \sum_{i=1}^m \ell(\mathbf{w}, \mathbf{x}_i, y_i)$$

which includes, e.g., all learning rules and approaches from Chapter 3